

Wine Data Set

This data is the result of a chemical analysis of wines grown in the same region in Italy but derived from 3 different cultivars. The analysis determined the quantities of 13 constituents found in each of the three types of wines.

Attribute Information

1. Type <-- (Label which has 3 classes)
2. Alcohol
3. Malic acid
4. Ash
5. Alcalinity of ash
6. Magnesium
7. Total phenols
8. Flavanoids
9. Nonflavanoid phenols
10. Proanthocyanins
11. Color intensity
12. Hue
13. OD280/OD315 of diluted wines
14. Proline

In [1]:

```
import numpy as np
import pandas as pd

from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score
```

In [2]:

```
df = pd.read_csv('/Users/rod/Documents/UCSD/COGS/COGS_118A/Project/Wine/wine.data.csv')
df.columns = ['Type', 'Alcohol', 'Malic Acid', 'Ash', 'Alcalinity of Ash', 'Magnesium', 'Total phenols', 'Flavanoids', 'Nonflavanoid phenols', 'Proanthocyanins', 'Color intensity', 'Hue', 'OD280/OD315', 'Proline']
```

In [3]:

```
print('df type: ' + str(type(df)))
print('df size: ' + str(df.shape))
df.head()
```

```
df type: <class 'pandas.core.frame.DataFrame'>
df size: (177, 14)
```

Out[3]:

	Type	Alcohol	Malic Acid	Ash	Alcalinity of Ash	Magnesium	Total phenols	Flavanoids	Nonflavanoic phenols
0	1	13.20	1.78	2.14	11.2	100	2.65	2.76	0.26
1	1	13.16	2.36	2.67	18.6	101	2.80	3.24	0.30
2	1	14.37	1.95	2.50	16.8	113	3.85	3.49	0.24
3	1	13.24	2.59	2.87	21.0	118	2.80	2.69	0.39
4	1	14.20	1.76	2.45	15.2	112	3.27	3.39	0.34

Oberving the Data Set

1. Checks for null values.
2. Sees how many classes/categories there are.
3. Counts the data points that belong to each category.

In [4]:

```
print('Number of NULL values in df: ' + str(df.isnull().sum().sum()))

uniqueClasses = df['Type'].unique()
print('Number of unique classes in df: ' + str(uniqueClasses.shape))
uniqueClasses = np.sort(uniqueClasses)

for i in uniqueClasses:
    print('Class ' + str(i) + ' count: ' + str((df['Type']==i).sum()))
```

```
Number of NULL values in df: 0
Number of unique classes in df: (3,)
Class 1 count: 58
Class 2 count: 71
Class 3 count: 48
```

Shuffle Data Randomly

1. Saves the first random shuffle of the original df.
2. Saves the second random shuffle of the original df.
3. Saves the third random shuffle of the original df.

In [5]:

```
df_shuffle1 = df.sample(frac=1)
#df_shuffle1.head()
```

In [6]:

```
df_shuffle2 = df.sample(frac=1)
#df_shuffle2.head()
```

In [7]:

```
df_shuffle3 = df.sample(frac=1)
#df_shuffle3.head()
```

$F(X) = Y$

Separates data into X and Y (labels) to set up the rest of the supervised learning algos in the [$F(X) = Y$] format.

1. Sets up $F(X1) = Y1$ from the first random shuffle of the original df.
2. Sets up $F(X2) = Y2$ from the second random shuffle of the original df.
3. Sets up $F(X3) = Y3$ from the third random shuffle of the original df.

In [8]:

```
df_array1 = np.array(df_shuffle1)           #Convert dataframe to array in order to
slice into X and Y.

X1 = df_array1[:,1:df_array1.shape[1]]      #Second Column to last column. Features
are all numerical.
Y1 = df_array1[:, 0]                        #First Column. This feature represents t
he classes or type of wine.
```

In [9]:

```
df_array2 = np.array(df_shuffle2)      #Convert dataframe to array in order to slice into X and Y.

X2 = df_array2[:,1:df_array2.shape[1]] #Second Column to last column. Features are all numerical.
Y2 = df_array2[:, 0]                   #First Column. This feature represents the classes or type of wine.
```

In [10]:

```
df_array3 = np.array(df_shuffle3)      #Convert dataframe to array in order to slice into X and Y.

X3 = df_array3[:,1:df_array3.shape[1]] #Second Column to last column. Features are all numerical.
Y3 = df_array3[:, 0]                   #First Column. This feature represents the classes or type of wine.
#print('X3 shape: ' + str(X3.shape))
#print('Y3 shape: ' + str(Y3.shape))
#print(X3[:, 0])
#print(Y3)
```

Functions Used For All Classifiers

1. partitionData
2. viewSplit
3. draw_heatmap_linear
4. bestValue **(EXTRA CREDIT)**
5. ViewConfusionMatrix
6. displayAccuracies

In [11]:

```
#X: Features of df.
#Y: Labels of df.
#percent: The percentage given to the training_validation set.
def partitionData(X, Y, percent):
    X_train_val = X[:int(percent*len(X))] # Get features from train + val set.
    Y_train_val = Y[:int(percent*len(Y))] # Get labels from train + val set.
    X_test      = X[int(percent*len(X)):] # Get features from test set.
    Y_test      = Y[int(percent*len(Y)):] # Get labels from test set.

    return X_train_val, Y_train_val, X_test, Y_test
```

In [12]:

```
#PURPOSE: Used to see the dimensions of the data after being partioned.
#Prints the shape of X_train_val.
#Prints the shape of Y_train_val.
#Prints the shape of X_test.
#Prints the shape of Y_test.
#Prints num of UNIQUE classes in Y_train_val.
#Prints the num of data points that belong to each class/category.
#Prints num of UNIQUE classes in Y_test.
def viewSplit(X_train_val, Y_train_val, X_test, Y_test):
    print('X_train_val shape: ' + str(X_train_val.shape))
    print('Y_train_val shape: ' + str(Y_train_val.shape))
    print('X_test: ' + str(X_test.shape))
    print('Y_test: ' + str(Y_test.shape))

    uniqueClasses = df['Type'].unique()
    print('Number of unique classes in df: ' + str(uniqueClasses.shape))
    uniqueClasses = np.sort(uniqueClasses)

    uniqueClasses_Y_train_val = np.unique(Y_train_val)
    print('Number of unique classes in Y_train_val: ' + str(uniqueClasses_Y_train_val.shape))
    for i in uniqueClasses:
        print('Class ' + str(i) + ' count: ' + str((Y_train_val[:,]==i).sum()))
    uniqueClasses_Y_test = np.unique(Y_test)
    print('Number of unique classes in Y_test: ' + str(uniqueClasses_Y_test.shape))
```

In [13]:

```
import seaborn as sns
import matplotlib.pyplot as plt

#PURPOSE: Draw heatmaps for result of grid search and find best C for validation set.
def draw_heatmap_linear(acc, acc_desc, C_list):
    plt.figure(figsize = (2,4))
    ax = sns.heatmap(acc, annot=True, fmt='.3f', yticklabels=C_list, xticklabels=[])
    ax.collections[0].colorbar.set_label("accuracy")
    ax.set(ylabel='$C$')
    plt.title(acc_desc + ' w.r.t $C$')
    sns.set_style("whitegrid", {'axes.grid' : False})
    plt.show()
```

In [14]:

```
#PURPOSE: Searches for the highest value in accuracyValidation, then uses the index of the highest value
# to find what value in the list caused this.
def bestValue(accuracyValidation, valueList):
    max_value_of_accV = np.max(accuracyValidation)
    max_index_of_accV = np.argmax(accuracyValidation)
    print('Largest value in accuracyValidation is ' + str(max_value_of_accV) + '
from index ' + str(max_index_of_accV) + '.')
    return valueList[max_index_of_accV]
```

In [15]:

```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

def ViewConfusionMatrix(Y_test, pred):
    print('Original labels:\n' + str(Y_test))
    print('Original Labels or Y_test shape: ' + str(Y_test.shape))
    print('Predicted labels:\n' + str(pred))

    #Note that the shape of the confusion matrix is not based on the shape of the Y_test or pred, but instead on
    #how many unique classes were inside of these.
    print('\nTest Accuracy Score: ' + str(accuracy_score(Y_test, pred)))
    print(classification_report(Y_test, pred))
    confusionMatrix = confusion_matrix(Y_test, pred)
    print('Confusion Matrix shape: ' + str(confusionMatrix.shape))
    print(confusionMatrix) #Remove because it takes up to much space.....
```

In [106]:

```
import matplotlib.pyplot as plt; plt.rcdefaults()
import matplotlib.pyplot as plt

def displayAccuracies(stringClfName, stringDataName, acc80_20, acc50_50, acc20_80):

    objects = ('80%train,20%test', '50%train,50%test', '20%train,80%test')
    y_pos = np.arange(len(objects))
    performance = [acc80_20, acc50_50, acc20_80]

    plt.bar(y_pos, performance, align='center', alpha=0.5)
    plt.xticks(y_pos, objects)
    plt.ylabel('Accuracy')
    #plt.title('Random Forest\'s Test Accuracies for 3 Partitions')
    plt.title(str(stringClfName) + '\\'s Test Accuracies of 3 Partitions on ' + str(stringDataName))
    plt.grid() #new
    plt.show()
```

In [17]:

```
#PURPOSE to print out the accuracies of 3 trials for each of the 3 partitions.
def printAccuracies(stringClgName, list_80_20, list_50_50, list20_80):

    print('Accuracy of ' + str(stringClgName) + '\s 3 trials on (80% train, 20%
test) partition :' + str(list_80_20))
    accuracyAverage_80_20 = np.mean(list_80_20)
    print('Mean Accuracy of ' + str(stringClgName) + ' on (80% train, 20% test) p
artition: ' + str(accuracyAverage_80_20))

    print('\nAccuracy of ' + str(stringClgName) + '\s 3 trials on (50% train, 5
0% test) partition :' + str(list_50_50))
    accuracyAverage_50_50 = np.mean(list_50_50)
    print('Mean Accuracy of ' + str(stringClgName) + ' on (50% train, 50% test)
partition: ' + str(accuracyAverage_50_50))

    print('\nAccuracy of ' + str(stringClgName) + '\s 3 trials on (20% train, 8
0% test) partition:' + str(list20_80))
    accuracyAverage_20_80 = np.mean(list20_80)
    print('Mean Accuracy of ' + str(stringClgName) + ' on (20% train, 80% test)
partition: ' + str(accuracyAverage_20_80))
```

Global Variables

CV: The number of folds that happens in the cross validation produced by GridSearchCV.

In [18]:

```
#Recommend running final test with CV=10, but during staging use CV=3.
CV = 10
```

Support Vector Machine (SVM)

Its a supervised machine learning algorithm which can be used for both classification or regression problems. But it is usually used for classification. Given 2 or more labeled classes of data, it acts as a discriminative classifier, formally defined by an optimal hyperplane that separates all the classes.

In [19]:

```
#GLOBAL VARIABLES FOR SVM
C_list = [10**-5, 10**-4, 10**-3, 10**-2, 10**-1, 1, 10, 100]
SVM_accuracyTestList_80_20 = []
SVM_accuracyTestList_50_50 = []
SVM_accuracyTestList_20_80 = []
```

In [20]:

```
from sklearn import svm

#C_list: C hyperparameter.
#cv: Number of folds when doing cross validation.
def svmTrainValidation(X_train_val, Y_train_val, C_list, CV):

    svm_classifier = svm.SVC(kernel = 'linear')

    parameters = {'C':C_list}

    SVM_clfGridSearch = GridSearchCV(svm_classifier, param_grid=parameters, cv=CV, return_train_score=True)
    SVM_clfGridSearch.fit(X_train_val, Y_train_val)

    #accuracyTrain = clfGridSearch.cv_results_['mean_train_score']
    #accuracyValidation = clfGridSearch.cv_results_['mean_test_score']
    return SVM_clfGridSearch
```

SVM on (80% train, 20% test)

Splits the 3 shuffled datasets into 2 parts:

1. (80% of all the data points) ---> Training set + Validation Set.
2. (20% of all the data points) ---> Test set.

In [21]:

```
X1_train_val, Y1_train_val, X1_test, Y1_test = partitionData(X1, Y1, 0.8)
X2_train_val, Y2_train_val, X2_test, Y2_test = partitionData(X2, Y2, 0.8)
X3_train_val, Y3_train_val, X3_test, Y3_test = partitionData(X3, Y3, 0.8)
```

1st Run)

First run uses the variables; X1_train_val, Y1_train_val, X1_test, Y1_test



In [22]:

```
SVM_clfGridSearch = svmTrainValidation(X1_train_val, Y1_train_val, C_list, CV)
accuracyTrain = SVM_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = SVM_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', C_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', C_list)
```



```
[ 0.68398668  0.69186128  0.78875117  0.93933333  0.97872579  1.
 1.
 1.
 ]
[ 0.68085106  0.68794326  0.77304965  0.91489362  0.96453901  0.9574
4681
 0.95035461  0.95035461]
```



In [23]:

```
#EXTRA CREDIT
#Instead of creating my own implementation to pick out the best hyper.
#I could have used this to find the best hyperparameter.
print(SVM_clfGridSearch.best_params_)

#Use the best C to calculate the test accuracy.
best_C = bestValue(accuracyValidation, C_list)
print('Best C: ' + str(best_C))

optimalClassifier = svm.SVC(kernel = 'linear', C = best_C).fit(X1_train_val, Y1_train_val)
pred = optimalClassifier.predict(X1_test)
# correct = [(a==b) for (a,b) in zip(pred,Y1_test)]
# test_acc = sum(correct) * 1.0 / len(correct)
# print('Test Accuracy Score: ' + str(test_acc))

#accuracy(ORIGINAL_VALUES, PREDICTED_VALUES)
accuracyTest = accuracy_score(Y1_test, pred)
SVM_accuracyTestList_80_20.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

```
{'C': 0.1}
```

Largest value in accuracyValidation is 0.964539007092 from index 4.

Best C: 0.1

Test Accuracy Score: 0.972222222222

2nd Run)



Second run uses the variables; X2_train_val, Y2_train_val, X2_test, Y2_test

In [24]:

```
SVM_clfGridSearch = svmTrainValidation(X2_train_val, Y2_train_val, C_list, CV)
accuracyTrain = SVM_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = SVM_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain) #This is what shows up in the heat maps.
print(accuracyValidation) #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', C_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', C_list)
```

```
[ 0.69974585  0.72334506  0.81635927  0.95272756  0.97791993  0.9992
1.
1.
]
[ 0.66666667  0.69503546  0.77304965  0.90780142  0.95035461  0.9574
4681
0.95035461  0.95035461]
```



In [25]:

```
best_C = bestValue(accuracyValidation, C_list)
print('Best C: ' + str(best_C))

optimalClassifier = svm.SVC(kernel = 'linear', C = best_C).fit(X2_train_val, Y2_train_val)
pred = optimalClassifier.predict(X2_test)

accuracyTest = accuracy_score(Y2_test, pred)
SVM_accuracyTestList_80_20.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.957446808511 from index 5.

Best C: 1

Test Accuracy Score: 0.972222222222

3rd Run)



Third run uses the variables; X3_train_val, Y3_train_val, X3_test, Y3_test

In [26]:

```
SVM_clfGridSearch = svmTrainValidation(X3_train_val, Y3_train_val, C_list, CV)
accuracyTrain = SVM_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = SVM_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)           #This is what shows up in the heat maps.
print(accuracyValidation)      #This is what shows up in the heat maps.
```

```
train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', C_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', C_list)
```

```
[ 0.71556132  0.71714257  0.8447841   0.94871629  0.9771513   0.9881
9425
  1.          1.          ]
[ 0.69503546  0.71631206  0.84397163  0.90780142  0.96453901  0.9503
5461
  0.95035461  0.95035461]
```



In [27]:

```
best_C = bestValue(accuracyValidation, C_list)
print('Best C: ' + str(best_C))

optimalClassifier = svm.SVC(kernel = 'linear', C = best_C).fit(X3_train_val, Y3_train_val)
pred = optimalClassifier.predict(X3_test)

accuracyTest = accuracy_score(Y3_test, pred)
SVM_accuracyTestList_80_20.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.964539007092 from index 4.

Best C: 0.1

Test Accuracy Score: 0.972222222222

Mean of SVM's Test Accuracies on (80% train, 20% test)

In [28]:

```
import statistics

print('SVM_accuracyTestList:' + str(SVM_accuracyTestList_80_20))
SVM_accuracyAverage_80_20 = statistics.mean(SVM_accuracyTestList_80_20)
print('SVM_accuracyTestList mean: ' + str(SVM_accuracyAverage_80_20))
```

SVM_accuracyTestList:[0.9722222222222221, 0.9722222222222221, 0.9722222222222221]

SVM_accuracyTestList mean: 0.972222222222

SVM on (50% train, 50% test)

Splits the 3 shuffled datasets into 2 parts:

1. (50% of all the data points) ---> Training set + Validation Set.
2. (50% of all the data points) ---> Test set.

In [29]:

```
X1_train_val, Y1_train_val, X1_test, Y1_test = partitionData(X1, Y1, 0.5)
X2_train_val, Y2_train_val, X2_test, Y2_test = partitionData(X2, Y2, 0.5)
X3_train_val, Y3_train_val, X3_test, Y3_test = partitionData(X3, Y3, 0.5)
```

1st Run)


First run uses the variables; X1_train_val, Y1_train_val, X1_test, Y1_test


In [30]:

```
SVM_clfGridSearch = svmTrainValidation(X1_train_val, Y1_train_val, C_list, CV)
accuracyTrain = SVM_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = SVM_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', C_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', C_list)
```

```
[ 0.6817953    0.68681113  0.7031577    0.93182725  0.97727561  1.
1.
1.
1.]
[ 0.68181818  0.68181818  0.69318182  0.86363636  0.95454545  0.9431
8182
0.94318182  0.94318182]
```





In [31]:

```
best_C = bestValue(accuracyValidation, C_list)
print('Best C: ' + str(best_C))

optimalClassifier = svm.SVC(kernel = 'linear', C = best_C).fit(X1_train_val, Y1_train_val)
pred = optimalClassifier.predict(X1_test)

accuracyTest = accuracy_score(Y1_test, pred)
SVM_accuracyTestList_50_50.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.954545454545 from index 4.

Best C: 0.1

Test Accuracy Score: 0.966292134831

2nd Run)



Second run uses the variables; X2_train_val, Y2_train_val, X2_test, Y2_test

In [32]:

```
SVM_clfGridSearch = svmTrainValidation(X2_train_val, Y2_train_val, C_list, CV)
accuracyTrain = SVM_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = SVM_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', C_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', C_list)
```

```
[ 0.69187678  0.69194088  0.81169753  0.93565052  0.98990266  1.
 1.
 1.
 ]
[ 0.68181818  0.67045455  0.77272727  0.875          0.95454545  0.9318
1818
 0.93181818  0.93181818]
```



In [33]:

```
best_C = bestValue(accuracyValidation, C_list)
print('Best C: ' + str(best_C))

optimalClassifier = svm.SVC(kernel = 'linear', C = best_C).fit(X2_train_val, Y2_train_val)
pred = optimalClassifier.predict(X2_test)

accuracyTest = accuracy_score(Y2_test, pred)
SVM_accuracyTestList_50_50.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.954545454545 from index 4.

Best C: 0.1

Test Accuracy Score: 0.955056179775

3rd Run)



Third run uses the variables; X3_train_val, Y3_train_val, X3_test, Y3_test

In [34]:

```
SVM_clfGridSearch = svmTrainValidation(X3_train_val, Y3_train_val, C_list, CV)
accuracyTrain = SVM_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = SVM_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', C_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', C_list)
```

```
[ 0.64404256  0.6464476   0.75249178  0.92035102  0.99746795  1.
 1.
 1.
 ]
[ 0.64772727  0.64772727  0.70454545  0.86363636  0.96590909  0.9431
8182
 0.94318182  0.94318182]
```

train accuracy w.r.t C val accuracy w.r.t C 

In [35]:

```
best_C = bestValue(accuracyValidation, C_list)
print('Best C: ' + str(best_C))

optimalClassifier = svm.SVC(kernel = 'linear', C = best_C).fit(X3_train_val, Y3_train_val)
pred = optimalClassifier.predict(X3_test)

accuracyTest = accuracy_score(Y3_test, pred)
SVM_accuracyTestList_50_50.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.965909090909 from index 4.

Best C: 0.1

Test Accuracy Score: 0.966292134831

Mean of SVM's Test Accuracies on (50% train, 50% test)

In [36]:

```
print('SVM_accuracyTestList:' + str(SVM_accuracyTestList_50_50))
SVM_accuracyAverage_50_50 = statistics.mean(SVM_accuracyTestList_50_50)
print('SVM_accuracyTestList mean: ' + str(SVM_accuracyAverage_50_50))
```

SVM_accuracyTestList:[0.9662921348314607, 0.9550561797752809, 0.9662921348314607]

SVM_accuracyTestList mean: 0.962546816479

SVM on (20% train, 80% test)

Splits the 3 shuffled datasets into 2 parts:

1. (20% of all the data points) ---> Training set + Validation Set.
2. (80% of all the data points) ---> Test set.

In [37]:

```
X1_train_val, Y1_train_val, X1_test, Y1_test = partitionData(X1, Y1, 0.2)
X2_train_val, Y2_train_val, X2_test, Y2_test = partitionData(X2, Y2, 0.2)
X3_train_val, Y3_train_val, X3_test, Y3_test = partitionData(X3, Y3, 0.5)
```

1st Run)

First run uses the variables; X1_train_val, Y1_train_val, X1_test, Y1_test


In [38]:


```
SVM_clfGridSearch = svmTrainValidation(X1_train_val, Y1_train_val, C_list, CV)
accuracyTrain = SVM_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = SVM_clfGridSearch.cv_results_['mean_test_score']

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', C_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', C_list)
```

/Users/rod/anaconda3/lib/python3.6/site-packages/sklearn/model_selection/_split.py:605: Warning: The least populated class in y has only 7 members, which is too few. The minimum number of members in any class cannot be less than n_splits=10.

% (min_groups, self.n_splits)), Warning)





In [39]:

```
best_C = bestValue(accuracyValidation, C_list)
print('Best C: ' + str(best_C))

optimalClassifier = svm.SVC(kernel = 'linear', C = best_C).fit(X1_train_val, Y1_train_val)
pred = optimalClassifier.predict(X1_test)

accuracyTest = accuracy_score(Y1_test, pred)
SVM_accuracyTestList_20_80.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.857142857143 from index 4.

Best C: 0.1

Test Accuracy Score: 0.859154929577



2nd Run)

Second run uses the variables; X2_train_val, Y2_train_val, X2_test, Y2_test

In [40]:

```
SVM_clfGridSearch = svmTrainValidation(X2_train_val, Y2_train_val, C_list, CV)
accuracyTrain = SVM_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = SVM_clfGridSearch.cv_results_['mean_test_score']

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', C_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', C_list)
```



In [41]:

```
best_C = bestValue(accuracyValidation, C_list)
print('Best C: ' + str(best_C))

optimalClassifier = svm.SVC(kernel = 'linear', C = best_C).fit(X2_train_val, Y2_train_val)
pred = optimalClassifier.predict(X2_test)

accuracyTest = accuracy_score(Y2_test, pred)
SVM_accuracyTestList_20_80.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.857142857143 from index 4.
Best C: 0.1
Test Accuracy Score: 0.880281690141

3rd Run)



Third run uses the variables; X3_train_val, Y3_train_val, X3_test, Y3_test

In [42]:

```
SVM_clfGridSearch = svmTrainValidation(X3_train_val, Y3_train_val, C_list, CV)
accuracyTrain = SVM_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = SVM_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)           #This is what shows up in the heat maps.
print(accuracyValidation)      #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', C_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', C_list)
```

```
[ 0.64404256  0.6464476   0.75249178  0.92035102  0.99746795  1.
 1.
 1.
      ]
[ 0.64772727  0.64772727  0.70454545  0.86363636  0.96590909  0.9431
8182
 0.94318182  0.94318182]
```

train accuracy w.r.t C val accuracy w.r.t C 

In [43]:

```
best_C = bestValue(accuracyValidation, C_list)
print('Best C: ' + str(best_C))

optimalClassifier = svm.SVC(kernel = 'linear', C = best_C).fit(X3_train_val, Y3_train_val)
pred = optimalClassifier.predict(X3_test)

accuracyTest = accuracy_score(Y3_test, pred)
SVM_accuracyTestList_20_80.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.965909090909 from index 4.

Best C: 0.1

Test Accuracy Score: 0.966292134831

Mean of SVM's Test Accuracies on (20% train, 80% test)

In [44]:

```
print('SVM_accuracyTestList:' + str(SVM_accuracyTestList_20_80))
SVM_accuracyAverage_20_80 = statistics.mean(SVM_accuracyTestList_20_80)
print('SVM_accuracyTestList mean: ' + str(SVM_accuracyAverage_20_80))
```

SVM_accuracyTestList:[0.85915492957746475, 0.88028169014084512, 0.9662921348314607]

SVM_accuracyTestList mean: 0.90190958485

Results of SVM

In [45]:

```
print('SVM_accuracyTestList (80% train, 20% test) partition mean: ' + str(SVM_accuracyAverage_80_20))
print('SVM_accuracyTestList (50% train, 50% test) partition mean: ' + str(SVM_accuracyAverage_50_50))
print('SVM_accuracyTestList (20% train, 80% test) partition mean: ' + str(SVM_accuracyAverage_20_80))
```

SVM_accuracyTestList (80% train, 20% test) partition mean: 0.9722222222

SVM_accuracyTestList (50% train, 50% test) partition mean: 0.962546816479

SVM_accuracyTestList (20% train, 80% test) partition mean: 0.90190958485

In [107]:


```
displayAccuracies('SVM', 'Wine Data', SVM_accuracyAverage_80_20, SVM_accuracyAverage_50_50, SVM_accuracyAverage_20_80)

printAccuracies('SVM', SVM_accuracyTestList_80_20, SVM_accuracyTestList_50_50, SVM_accuracyTestList_20_80)
# print('Accuracy of SVM\'s 3 trials on (80% train, 20% test) partition : ' + str(SVM_accuracyTestList_80_20))
# SVM_accuracyAverage_80_20 = statistics.mean(SVM_accuracyTestList_80_20)
# print('Mean Accuracy of SVM on (80% train, 20% test) partition: ' + str(SVM_accuracyAverage_80_20))

# print('\nAccuracy of SVM\'s 3 trials on (50% train, 50% test) partition : ' + str(SVM_accuracyTestList_50_50))
# SVM_accuracyAverage_50_50 = statistics.mean(SVM_accuracyTestList_50_50)
# print('Mean Accuracy of SVM on (50% train, 50% test) partition: ' + str(SVM_accuracyAverage_50_50))

# print('\nAccuracy of SVM\'s 3 trials on (20% train, 80% test) partition: ' + str(SVM_accuracyTestList_20_80))
# SVM_accuracyAverage_20_80 = statistics.mean(SVM_accuracyTestList_20_80)
# print('Mean Accuracy of SVM on (20% train, 80% test) partition: ' + str(SVM_accuracyAverage_20_80))
```

SVM's Test Accuracies of 3 Partitions on Wine Data



Accuracy of SVM's 3 trials on (80% train, 20% test) partition :[0.972222222222221, 0.972222222222221, 0.972222222222221]
 Mean Accuracy of SVM on (80% train, 20% test) partition: 0.972222222222

Accuracy of SVM's 3 trials on (50% train, 50% test) partition :[0.9662921348314607, 0.9550561797752809, 0.9662921348314607]
 Mean Accuracy of SVM on (50% train, 50% test) partition: 0.962546816479

Accuracy of SVM's 3 trials on (20% train, 80% test) partition:[0.85915492957746475, 0.88028169014084512, 0.9662921348314607]
 Mean Accuracy of SVM on (20% train, 80% test) partition: 0.90190958485

Decision Tree

Decision Trees (DTs) are a non-parametric supervised learning method used for classification and regression. The goal is to create a model that predicts the value of a target variable by learning simple decision rules inferred from the data features.

In [47]:

```
#GLOBAL VARIABLES FOR Decision Tree
#D_list = np.asarray([1,2,3,4,5])
D_list = [1,2,3,4,5]
DT_accuracyTestList_80_20 = []
DT_accuracyTestList_50_50 = []
DT_accuracyTestList_20_80 = []
```

In [48]:

```
from sklearn import tree

def decisionTreeTrainValidation(X_train_val, Y_train_val, D_list, CV):

    DT_classifier = tree.DecisionTreeClassifier(criterion='entropy')

    parameters = {'max_depth': D_list}

    DT_clfGridSearch = GridSearchCV(DT_classifier, param_grid=parameters, cv=CV,
    return_train_score=True)
    DT_clfGridSearch.fit(X_train_val, Y_train_val)

    return DT_clfGridSearch
```

Decision Tree on (80% train, 20% test)

Splits the 3 shuffled datasets into 2 parts:

1. (80% of all the data points) ---> Training set + Validation Set.
2. (20% of all the data points) ---> Test set.

In [49]:

```
X1_train_val, Y1_train_val, X1_test, Y1_test = partitionData(X1, Y1, 0.8)
X2_train_val, Y2_train_val, X2_test, Y2_test = partitionData(X2, Y2, 0.8)
X3_train_val, Y3_train_val, X3_test, Y3_test = partitionData(X3, Y3, 0.8)
```

1st Run)



First run uses the variables; X1_train_val, Y1_train_val, X1_test, Y1_test

In [50]:

```
DT_clfGridSearch = decisionTreeTrainValidation(X1_train_val, Y1_train_val, D_list, CV)
accuracyTrain = DT_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = DT_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', D_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', D_list)
```

```
[ 0.63989997  0.93207172  0.98974424  0.9992126   1.          ]
[ 0.58865248  0.85106383  0.92198582  0.90780142  0.90780142]
```



In [51]:

```
best_D = bestValue(accuracyValidation, D_list)
print('Best D: ' + str(best_D))

optimalClassifier = tree.DecisionTreeClassifier(criterion='entropy', max_depth=best_D).fit(X1_train_val, Y1_train_val)
pred = optimalClassifier.predict(X1_test)

accuracyTest = accuracy_score(Y1_test, pred)
DT_accuracyTestList_80_20.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.921985815603 from index 2.

Best D: 3

Test Accuracy Score: 1.0

2nd Run)



Second run uses the variables; X2_train_val, Y2_train_val, X2_test, Y2_test

In [52]:

```
DT_clfGridSearch = decisionTreeTrainValidation(X2_train_val, Y2_train_val, D_list, CV)
accuracyTrain = DT_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = DT_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', D_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', D_list)
```

```
[ 0.64394344  0.96524523  0.99131914  0.99921875  1.          ]
[ 0.60992908  0.89361702  0.92907801  0.92198582  0.93617021]
```



In [53]:

```
#Use the best C to calculate the test accuracy.
best_D = bestValue(accuracyValidation, D_list)
print('Best D: ' + str(best_D))

optimalClassifier = tree.DecisionTreeClassifier(criterion='entropy', max_depth=best_D).fit(X2_train_val, Y2_train_val)
pred = optimalClassifier.predict(X2_test)

accuracyTest = accuracy_score(Y2_test, pred)
DT_accuracyTestList_80_20.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.936170212766 from index 4.

Best D: 5

Test Accuracy Score: 0.916666666667

3rd Run)



Third run uses the variables; X3_train_val, Y3_train_val, X3_test, Y3_test

In [54]:

```
DT_clfGridSearch = decisionTreeTrainValidation(X3_train_val, Y3_train_val, D_list, CV)
accuracyTrain = DT_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = DT_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)           #This is what shows up in the heat maps.
print(accuracyValidation)      #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', D_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', D_list)
```

```
[ 0.61387406  0.96455835  0.9921255   1.          1.          ]
[ 0.59574468  0.92198582  0.92907801  0.94326241  0.92198582]
```



In [55]:

```
#Use the best C to calculate the test accuracy.
best_D = bestValue(accuracyValidation, D_list)
print('Best D: ' + str(best_D))

optimalClassifier = tree.DecisionTreeClassifier(criterion='entropy', max_depth=best_D).fit(X3_train_val, Y3_train_val)
pred = optimalClassifier.predict(X3_test)

accuracyTest = accuracy_score(Y3_test, pred)
DT_accuracyTestList_80_20.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.943262411348 from index 3.

Best D: 4

Test Accuracy Score: 1.0

Mean of DT's Test Accuracies on (80% train, 20% test)

In [56]:

```
print('DT_accuracyTestList:' + str(DT_accuracyTestList_80_20))
DT_accuracyAverage_80_20 = statistics.mean(DT_accuracyTestList_80_20)
print('DT_accuracyTestList mean: ' + str(DT_accuracyAverage_80_20))
```

DT_accuracyTestList:[1.0, 0.91666666666666663, 1.0]

DT_accuracyTestList mean: 0.972222222222

Decision Tree on (50% train, 50% test)

Splits the 3 shuffled datasets into 2 parts:

1. (50% of all the data points) ---> Training set + Validation Set.
2. (50% of all the data points) ---> Test set.

In [57]:

```
X1_train_val, Y1_train_val, X1_test, Y1_test = partitionData(X1, Y1, 0.5)
X2_train_val, Y2_train_val, X2_test, Y2_test = partitionData(X2, Y2, 0.5)
X3_train_val, Y3_train_val, X3_test, Y3_test = partitionData(X3, Y3, 0.5)
```

1st Run)


First run uses the variables; X1_train_val, Y1_train_val, X1_test, Y1_test


In [58]:

```
DT_clfGridSearch = decisionTreeTrainValidation(X1_train_val, Y1_train_val, D_list, CV)
accuracyTrain = DT_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = DT_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)           #This is what shows up in the heat maps.
print(accuracyValidation)      #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', D_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', D_list)
```

```
[ 0.69951804  0.94835034  0.99873418  0.99873418  1.          ]
[ 0.63636364  0.88636364  0.93181818  0.94318182  0.92045455]
```





In [59]:

```
best_D = bestValue(accuracyValidation, D_list)
print('Best D: ' + str(best_D))

optimalClassifier = tree.DecisionTreeClassifier(criterion='entropy', max_depth=best_D).fit(X1_train_val, Y1_train_val)
pred = optimalClassifier.predict(X1_test)

accuracyTest = accuracy_score(Y1_test, pred)
DT_accuracyTestList_50_50.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.943181818182 from index 3.

Best D: 4

Test Accuracy Score: 0.898876404494

2nd Run)

Second run uses the variables; X2_train_val, Y2_train_val, X2_test, Y2_test



In [60]:

```
DT_clfGridSearch = decisionTreeTrainValidation(X2_train_val, Y2_train_val, D_list, CV)
accuracyTrain = DT_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = DT_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', D_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', D_list)
```



```
[ 0.64273742  0.97982194  0.99370133  1.          1.          ]
[ 0.59090909  0.90909091  0.94318182  0.90909091  0.93181818]
```



In [61]:

```
best_D = bestValue(accuracyValidation, D_list)
print('Best D: ' + str(best_D))

optimalClassifier = tree.DecisionTreeClassifier(criterion='entropy', max_depth=best_D).fit(X2_train_val, Y2_train_val)
pred = optimalClassifier.predict(X2_test)

accuracyTest = accuracy_score(Y2_test, pred)
DT_accuracyTestList_50_50.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.943181818182 from index 2.

Best D: 3

Test Accuracy Score: 0.876404494382

3rd Run)

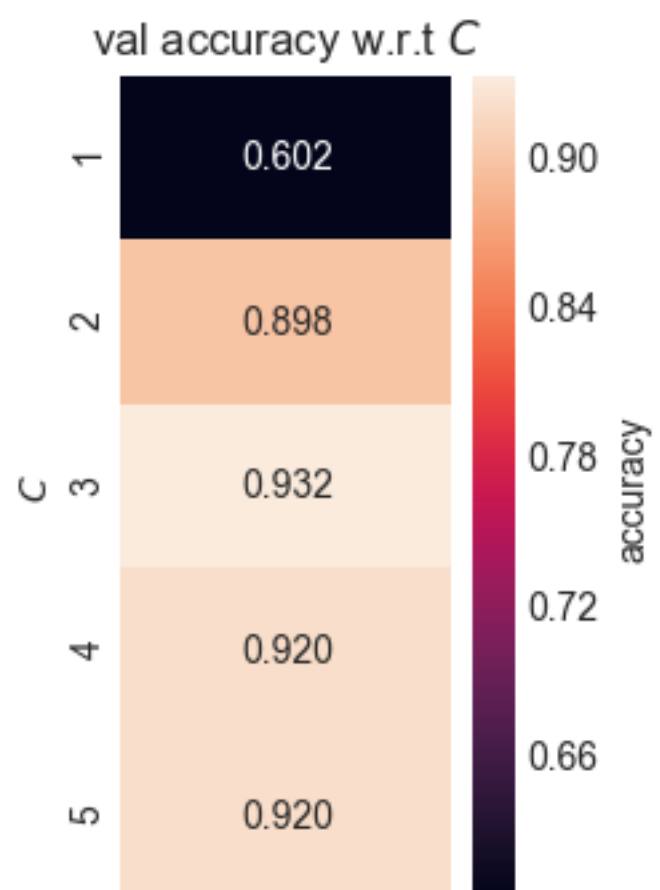

Third run uses the variables; X3_train_val, Y3_train_val, X3_test, Y3_test

In [62]:

```
DT_clfGridSearch = decisionTreeTrainValidation(X3_train_val, Y3_train_val, D_list, CV)
accuracyTrain = DT_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = DT_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', D_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', D_list)
```

```
[ 0.67308179  0.95208136  0.98991928  1.          1.          ]
[ 0.60227273  0.89772727  0.93181818  0.92045455  0.92045455]
```



In [63]:

```
best_D = bestValue(accuracyValidation, D_list)
print('Best D: ' + str(best_D))

optimalClassifier = tree.DecisionTreeClassifier(criterion='entropy', max_depth=best_D).fit(X3_train_val, Y3_train_val)
pred = optimalClassifier.predict(X3_test)

accuracyTest = accuracy_score(Y3_test, pred)
DT_accuracyTestList_50_50.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.931818181818 from index 2.

Best D: 3

Test Accuracy Score: 0.955056179775

Mean of DT's Test Accuracies on (50% train, 50% test)

In [64]:

```
print('DT_accuracyTestList:' + str(DT_accuracyTestList_50_50))
DT_accuracyAverage_50_50 = statistics.mean(DT_accuracyTestList_50_50)
print('DT_accuracyTestList mean: ' + str(DT_accuracyAverage_50_50))
```

DT_accuracyTestList:[0.898876404494382, 0.8764044943820225, 0.9550561797752809]

DT_accuracyTestList mean: 0.910112359551

Decision Tree on (20% train, 80% test)

Splits the 3 shuffled datasets into 2 parts:

1. (20% of all the data points) ---> Training set + Validation Set.
2. (80% of all the data points) ---> Test set.

In [65]:

```
X1_train_val, Y1_train_val, X1_test, Y1_test = partitionData(X1, Y1, 0.2)
X2_train_val, Y2_train_val, X2_test, Y2_test = partitionData(X2, Y2, 0.2)
X3_train_val, Y3_train_val, X3_test, Y3_test = partitionData(X3, Y3, 0.2)
```

1st Run)

First run uses the variables; X1_train_val, Y1_train_val, X1_test, Y1_test

In [66]:


```
DT_clfGridSearch = decisionTreeTrainValidation(X1_train_val, Y1_train_val, D_list, CV)
accuracyTrain = DT_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = DT_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.


train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', D_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', D_list)
```

```
[ 0.77475688  1.          1.          1.          1.          ]
[ 0.71428571  0.85714286  0.88571429  0.85714286  0.82857143]
```

/Users/rod/anaconda3/lib/python3.6/site-packages/sklearn/model_selection/_split.py:605: Warning: The least populated class in y has only 7 members, which is too few. The minimum number of members in any class cannot be less than n_splits=10.

```
% (min_groups, self.n_splits)), Warning)
```





In [67]:

```
best_D = bestValue(accuracyValidation, D_list)
print('Best D: ' + str(best_D))

optimalClassifier = tree.DecisionTreeClassifier(criterion='entropy', max_depth=best_D).fit(X1_train_val, Y1_train_val)
pred = optimalClassifier.predict(X1_test)

accuracyTest = accuracy_score(Y1_test, pred)
DT_accuracyTestList_20_80.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.885714285714 from index 2.

Best D: 3

Test Accuracy Score: 0.830985915493

2nd Run)



Second run uses the variables; X2_train_val, Y2_train_val, X2_test, Y2_test

In [68]:

```
DT_clfGridSearch = decisionTreeTrainValidation(X2_train_val, Y2_train_val, D_list, CV)
accuracyTrain = DT_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = DT_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)    #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', D_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', D_list)
```

```
[ 0.67635081  0.98386425  1.          1.          1.          ]
[ 0.51428571  0.68571429  0.71428571  0.77142857  0.74285714]
```



In [69]:

```
best_D = bestValue(accuracyValidation, D_list)
print('Best D: ' + str(best_D))

optimalClassifier = tree.DecisionTreeClassifier(criterion='entropy', max_depth=best_D).fit(X2_train_val, Y2_train_val)
pred = optimalClassifier.predict(X2_test)

accuracyTest = accuracy_score(Y2_test, pred)
DT_accuracyTestList_20_80.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.771428571429 from index 3.

Best D: 4

Test Accuracy Score: 0.901408450704

3rd Run)

Third run uses the variables; X3_train_val, Y3_train_val, X3_test, Y3_test

In [70]:



```
DT_clfGridSearch = decisionTreeTrainValidation(X3_train_val, Y3_train_val, D_list, CV)
accuracyTrain = DT_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = DT_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)           #This is what shows up in the heat maps.
print(accuracyValidation)      #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', D_list)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', D_list)
```

```
[ 0.7459708  0.96823375  1.          1.          1.          ]
[ 0.57142857  0.82857143  0.82857143  0.82857143  0.74285714]
```

/Users/rod/anaconda3/lib/python3.6/site-packages/sklearn/model_selection/_split.py:605: Warning: The least populated class in y has only 7 members, which is too few. The minimum number of members in any class cannot be less than n_splits=10.

```
% (min_groups, self.n_splits)), Warning)
```



In [71]:

```
best_D = bestValue(accuracyValidation, D_list)
print('Best D: ' + str(best_D))

optimalClassifier = tree.DecisionTreeClassifier(criterion='entropy', max_depth=best_D).fit(X3_train_val, Y3_train_val)
pred = optimalClassifier.predict(X3_test)

accuracyTest = accuracy_score(Y3_test, pred)
DT_accuracyTestList_20_80.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.828571428571 from index 1.

Best D: 2

Test Accuracy Score: 0.852112676056

Mean of DT's Test Accuracies on (20% train, 80% test)

In [72]:

```
print('DT_accuracyTestList:' + str(DT_accuracyTestList_20_80))
DT_accuracyAverage_20_80 = statistics.mean(DT_accuracyTestList_20_80)
print('DT_accuracyTestList mean: ' + str(DT_accuracyAverage_20_80))
```

DT_accuracyTestList:[0.83098591549295775, 0.90140845070422537, 0.852112676056338]

DT_accuracyTestList mean: 0.861502347418

Results of Decision Tree

In [73]:

```
print('DT_accuracyTestList (80% train, 20% test) partition mean: ' + str(DT_accuracyAverage_80_20))
print('DT_accuracyTestList (50% train, 50% test) partition mean: ' + str(DT_accuracyAverage_50_50))
print('DT_accuracyTestList (20% train, 80% test) partition mean: ' + str(DT_accuracyAverage_20_80))
```


DT_accuracyTestList (80% train, 20% test) partition mean: 0.972222222222

DT_accuracyTestList (50% train, 50% test) partition mean: 0.910112359551

DT_accuracyTestList (20% train, 80% test) partition mean: 0.861502347418

In [108]:

```
displayAccuracies('DT', 'Wine Data', DT_accuracyAverage_80_20, DT_accuracyAverage_50_50, DT_accuracyAverage_20_80)
printAccuracies('DT', DT_accuracyTestList_80_20, DT_accuracyTestList_50_50, DT_accuracyTestList_20_80)
```



Accuracy of DT's 3 trials on (80% train, 20% test) partition :[1.0, 0.9166666666666666, 1.0]
Mean Accuracy of DT on (80% train, 20% test) partition: 0.972222222222

Accuracy of DT's 3 trials on (50% train, 50% test) partition :[0.898876404494382, 0.8764044943820225, 0.9550561797752809]
Mean Accuracy of DT on (50% train, 50% test) partition: 0.910112359551

Accuracy of DT's 3 trials on (20% train, 80% test) partition:[0.83098591549295775, 0.90140845070422537, 0.852112676056338]
Mean Accuracy of DT on (20% train, 80% test) partition: 0.861502347418

Random Forest

A random forest is a meta estimator that fits a number of decision tree classifiers on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control over-fitting. The sub-sample size is always the same as the original input sample size but the samples are drawn with replacement if `bootstrap=True` (default).

In [75]:

```
#Global Variables For Random Forest
max_depth_List = [1,2,3,4,5]
RF_accuracyTestList_80_20 = []
RF_accuracyTestList_50_50 = []
RF_accuracyTestList_20_80 = []
```

In [76]:

```
from sklearn.ensemble import RandomForestClassifier

#max_depth_List: The chosen hyperparameter.
#cv: Number of folds when doing cross validation.
def randomForestTrainValidation(X_train_val, Y_train_val, max_depth_List, CV):

    #svm_classifier = svm.SVC(kernel = 'linear')
    RF_classifier = RandomForestClassifier()

    parameters = {'max_depth':max_depth_List}

    # param_grid = {
    #     'bootstrap': [True],
    #     'max_depth': [80, 90, 100, 110],
    #     'max_features': [2, 3],
    #     'min_samples_leaf': [3, 4, 5],
    #     'min_samples_split': [8, 10, 12],
    #     'n_estimators': [100, 200, 300, 1000]
    # }

    RF_clfGridSearch = GridSearchCV(RF_classifier, param_grid=parameters, cv=CV,
    return_train_score=True)
    RF_clfGridSearch.fit(X_train_val, Y_train_val)

    return RF_clfGridSearch
```

Random Forest on (80% train, 20% test)

Splits the 3 shuffled datasets into 2 parts:

1. (80% of all the data points) ---> Training set + Validation Set.
2. (20% of all the data points) ---> Test set.

In [77]:

```
X1_train_val, Y1_train_val, X1_test, Y1_test = partitionData(X1, Y1, 0.8)
X2_train_val, Y2_train_val, X2_test, Y2_test = partitionData(X2, Y2, 0.8)
X3_train_val, Y3_train_val, X3_test, Y3_test = partitionData(X3, Y3, 0.8)
```

1st Run)



First run uses the variables; X1_train_val, Y1_train_val, X1_test, Y1_test

In [78]:

```
RF_clfGridSearch = randomForestTrainValidation(X1_train_val, Y1_train_val, max_d
epth_List, CV)
accuracyTrain = RF_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = RF_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.
```

```
train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', max_depth_List)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', max_depth_List)
```

```
[ 0.91564771  0.96695734  0.9928877   0.9984251   0.99763135]
[ 0.89361702  0.93617021  0.94326241  0.93617021  0.95744681]
```



In [79]:

```
best_max_depth = bestValue(accuracyValidation, max_depth_List)
print('Best max_depth: ' + str(best_max_depth))

optimalClassifier = RandomForestClassifier(max_depth=best_max_depth).fit(X1_train_val, Y1_train_val)
pred = optimalClassifier.predict(X1_test)

accuracyTest = accuracy_score(Y1_test, pred)
RF_accuracyTestList_80_20.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.957446808511 from index 4.
Best max_depth: 5
Test Accuracy Score: 1.0

2nd Run)



Second run uses the variables; X2_train_val, Y2_train_val, X2_test, Y2_test

In [80]:

```
RF_clfGridSearch = randomForestTrainValidation(X2_train_val, Y2_train_val, max_depth_List, CV)
accuracyTrain = RF_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = RF_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)           #This is what shows up in the heat maps.
print(accuracyValidation)      #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', max_depth_List)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', max_depth_List)
```

```
[ 0.92354125  0.97870089  0.98973174  0.9984         1.          ]
[ 0.91489362  0.92198582  0.97163121  0.95035461  0.9787234  ]
```

In [81]:

```
best_max_depth = bestValue(accuracyValidation, max_depth_List)
print('Best max_depth: ' + str(best_max_depth))

optimalClassifier = RandomForestClassifier(max_depth=best_max_depth).fit(X2_train_val, Y2_train_val)
pred = optimalClassifier.predict(X2_test)

accuracyTest = accuracy_score(Y2_test, pred)
RF_accuracyTestList_80_20.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.978723404255 from index 4.

Best max_depth: 5

Test Accuracy Score: 1.0

3rd Run)



Third run uses the variables; X3_train_val, Y3_train_val, X3_test, Y3_test

In [82]:

```
RF_clfGridSearch = randomForestTrainValidation(X3_train_val, Y3_train_val, max_depth_List, CV)
accuracyTrain = RF_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = RF_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)           #This is what shows up in the heat maps.
print(accuracyValidation)      #This is what shows up in the heat maps.
```

```
train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', max_depth_List)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', max_depth_List)
```

```
[ 0.9464033    0.96927094   0.98818195   0.99530606   0.99920635]
[ 0.91489362   0.93617021   0.94326241   0.9787234    0.95744681]
```



In [83]:

```
best_max_depth = bestValue(accuracyValidation, max_depth_List)
print('Best max_depth: ' + str(best_max_depth))

optimalClassifier = RandomForestClassifier(max_depth=best_max_depth).fit(X3_train_val, Y3_train_val)
pred = optimalClassifier.predict(X3_test)

accuracyTest = accuracy_score(Y3_test, pred)
RF_accuracyTestList_80_20.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.978723404255 from index 3.

Best max_depth: 4

Test Accuracy Score: 1.0

Mean of RF's Test Accuracies on (80% train, 20% test)

In [84]:

```
print('RF_accuracyTestList:' + str(RF_accuracyTestList_80_20))
RF_accuracyAverage_80_20 = statistics.mean(RF_accuracyTestList_80_20)
print('RF_accuracyTestList mean: ' + str(RF_accuracyAverage_80_20))
```

RF_accuracyTestList:[1.0, 1.0, 1.0]

RF_accuracyTestList mean: 1.0

Random Forest on (50% train, 50% test)

Splits the 3 shuffled datasets into 2 parts:

1. (50% of all the data points) ---> Training set + Validation Set.
2. (50% of all the data points) ---> Test set.

In [85]:

```
X1_train_val, Y1_train_val, X1_test, Y1_test = partitionData(X1, Y1, 0.5)
X2_train_val, Y2_train_val, X2_test, Y2_test = partitionData(X2, Y2, 0.5)
X3_train_val, Y3_train_val, X3_test, Y3_test = partitionData(X3, Y3, 0.5)
```

1st Run)


First run uses the variables; X1_train_val, Y1_train_val, X1_test, Y1_test


In [86]:

```
RF_clfGridSearch = randomForestTrainValidation(X1_train_val, Y1_train_val, max_d
epth_List, CV)
accuracyTrain = RF_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = RF_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)           #This is what shows up in the heat maps.
print(accuracyValidation)      #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', max_depth_List)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', max_depth_List)
```

```
[ 0.92670068  0.98488801  0.99488677  0.9987013   0.9962013 ]
[ 0.89772727  0.95454545  0.98863636  0.97727273  0.93181818]
```





In [87]:

```
best_max_depth = bestValue(accuracyValidation, max_depth_List)
print('Best max_depth: ' + str(best_max_depth))

optimalClassifier = RandomForestClassifier(max_depth=best_max_depth).fit(X1_train_val, Y1_train_val)
pred = optimalClassifier.predict(X1_test)

accuracyTest = accuracy_score(Y1_test, pred)
RF_accuracyTestList_50_50.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.988636363636 from index 2.

Best max_depth: 3

Test Accuracy Score: 0.955056179775

2nd Run)



Second run uses the variables; X2_train_val, Y2_train_val, X2_test, Y2_test

In [88]:

```
RF_clfGridSearch = randomForestTrainValidation(X2_train_val, Y2_train_val, max_d
epth_List, CV)
accuracyTrain = RF_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = RF_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)    #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', max_depth_List)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', max_depth_List)
```

```
[ 0.94827873  0.98620014  1.          1.          0.99746795]
[ 0.89772727  0.94318182  0.96590909  0.98863636  0.96590909]
```



In [89]:

```
best_max_depth = bestValue(accuracyValidation, max_depth_List)
print('Best max_depth: ' + str(best_max_depth))

optimalClassifier = RandomForestClassifier(max_depth=best_max_depth).fit(X2_train_val, Y2_train_val)
pred = optimalClassifier.predict(X2_test)

accuracyTest = accuracy_score(Y2_test, pred)
RF_accuracyTestList_50_50.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.988636363636 from index 3.

Best max_depth: 4

Test Accuracy Score: 0.876404494382

3rd Run)



Third run uses the variables; X3_train_val, Y3_train_val, X3_test, Y3_test

In [90]:

```
RF_clfGridSearch = randomForestTrainValidation(X3_train_val, Y3_train_val, max_depth_List, CV)
accuracyTrain = RF_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = RF_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)           #This is what shows up in the heat maps.
print(accuracyValidation)      #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', max_depth_List)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', max_depth_List)
```

```
[ 0.90376832  0.97613485  0.99623338  0.99746795  1.          ]
[ 0.86363636  0.92045455  0.92045455  0.95454545  0.90909091]
```



In [91]:

```
best_max_depth = bestValue(accuracyValidation, max_depth_List)
print('Best max_depth: ' + str(best_max_depth))

optimalClassifier = RandomForestClassifier(max_depth=best_max_depth).fit(X3_train_val, Y3_train_val)
pred = optimalClassifier.predict(X3_test)

accuracyTest = accuracy_score(Y3_test, pred)
RF_accuracyTestList_50_50.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.954545454545 from index 3.

Best max_depth: 4

Test Accuracy Score: 0.966292134831

Mean of RF's Test Accuracies on (50% train, 50% test)

In [92]:

```
print('RF_accuracyTestList:' + str(RF_accuracyTestList_50_50))
RF_accuracyAverage_50_50 = statistics.mean(RF_accuracyTestList_50_50)
print('RF_accuracyTestList mean: ' + str(RF_accuracyAverage_50_50))
```

RF_accuracyTestList:[0.9550561797752809, 0.8764044943820225, 0.9662921348314607]

RF_accuracyTestList mean: 0.932584269663

Random Forest on (20% train, 80% test)

Splits the 3 shuffled datasets into 2 parts:

1. (20% of all the data points) ---> Training set + Validation Set.
2. (80% of all the data points) ---> Test set.

In [93]:

```
X1_train_val, Y1_train_val, X1_test, Y1_test = partitionData(X1, Y1, 0.2)
X2_train_val, Y2_train_val, X2_test, Y2_test = partitionData(X2, Y2, 0.2)
X3_train_val, Y3_train_val, X3_test, Y3_test = partitionData(X3, Y3, 0.2)
```

1st Run)

First run uses the variables; X1_train_val, Y1_train_val, X1_test, Y1_test

In [94]:


```
RF_clfGridSearch = randomForestTrainValidation(X1_train_val, Y1_train_val, max_depth_List, CV)
accuracyTrain = RF_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = RF_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.


train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', max_depth_List)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', max_depth_List)
```

/Users/rod/anaconda3/lib/python3.6/site-packages/sklearn/model_selection/_split.py:605: Warning: The least populated class in y has only 7 members, which is too few. The minimum number of members in any class cannot be less than n_splits=10.

% (min_groups, self.n_splits)), Warning)

```
[ 0.95514113  0.99354839  1.          0.99677419  1.          ]
[ 0.8         0.91428571  0.91428571  0.91428571  0.94285714]
```





In [95]:

```
best_max_depth = bestValue(accuracyValidation, max_depth_List)
print('Best max_depth: ' + str(best_max_depth))

optimalClassifier = RandomForestClassifier(max_depth=best_max_depth).fit(X1_train_val, Y1_train_val)
pred = optimalClassifier.predict(X1_test)

accuracyTest = accuracy_score(Y1_test, pred)
RF_accuracyTestList_20_80.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.942857142857 from index 4.

Best max_depth: 5

Test Accuracy Score: 0.894366197183

2nd Run)



Second run uses the variables; X2_train_val, Y2_train_val, X2_test, Y2_test

In [96]:

```
RF_clfGridSearch = randomForestTrainValidation(X2_train_val, Y2_train_val, max_d
epth_List, CV)
accuracyTrain = RF_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = RF_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', max_depth_List)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', max_depth_List)
```

```
[ 0.91188844  0.996875    1.         1.         0.99375    ]
[ 0.8         0.91428571  0.85714286  0.91428571  0.94285714]
```



In [97]:

```
best_max_depth = bestValue(accuracyValidation, max_depth_List)
print('Best max_depth: ' + str(best_max_depth))

optimalClassifier = RandomForestClassifier(max_depth=best_max_depth).fit(X2_train_val, Y2_train_val)
pred = optimalClassifier.predict(X2_test)

accuracyTest = accuracy_score(Y2_test, pred)
RF_accuracyTestList_20_80.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.942857142857 from index 4.

Best max_depth: 5

Test Accuracy Score: 0.887323943662

3rd Run)

Third run uses the variables; X3_train_val, Y3_train_val, X3_test, Y3_test

In [98]:



```
RF_clfGridSearch = randomForestTrainValidation(X3_train_val, Y3_train_val, max_depth_List, CV)
accuracyTrain = RF_clfGridSearch.cv_results_['mean_train_score']
accuracyValidation = RF_clfGridSearch.cv_results_['mean_test_score']
print(accuracyTrain)          #This is what shows up in the heat maps.
print(accuracyValidation)     #This is what shows up in the heat maps.

train_acc = (accuracyTrain).reshape(-1,1)
draw_heatmap_linear(train_acc, 'train accuracy', max_depth_List)
val_acc = (accuracyValidation).reshape(-1,1)
draw_heatmap_linear(val_acc, 'val accuracy', max_depth_List)
```

/Users/rod/anaconda3/lib/python3.6/site-packages/sklearn/model_selection/_split.py:605: Warning: The least populated class in y has only 7 members, which is too few. The minimum number of members in any class cannot be less than n_splits=10.

% (min_groups, self.n_splits)), Warning)

```
[ 0.83715481  0.98406647  1.          0.99071359  1.          ]
[ 0.71428571  0.77142857  0.91428571  0.91428571  0.8          ]
```

In [99]:

```
best_max_depth = bestValue(accuracyValidation, max_depth_List)
print('Best max_depth: ' + str(best_max_depth))

optimalClassifier = RandomForestClassifier(max_depth=best_max_depth).fit(X3_train_val, Y3_train_val)
pred = optimalClassifier.predict(X3_test)

accuracyTest = accuracy_score(Y3_test, pred)
RF_accuracyTestList_20_80.append(accuracyTest)
print('Test Accuracy Score: ' + str(accuracyTest))
```

Largest value in accuracyValidation is 0.914285714286 from index 2.
Best max_depth: 3
Test Accuracy Score: 0.887323943662

Mean of RF's Test Accuracies on (20% train, 80% test)

In [100]:

```
print('RT_accuracyTestList:' + str(RF_accuracyTestList_20_80))
RF_accuracyAverage_20_80 = statistics.mean(RF_accuracyTestList_20_80)
print('RT_accuracyTestList mean: ' + str(RF_accuracyAverage_20_80))
```

RT_accuracyTestList:[0.89436619718309862, 0.88732394366197187, 0.88732394366197187]
RT_accuracyTestList mean: 0.889671361502

Results of Random Forest


In [101]:

```
print('RT_accuracyTestList (80% train, 20% test) partition mean: ' + str(RF_accuracyAverage_80_20))
print('RT_accuracyTestList (50% train, 50% test) partition mean: ' + str(RF_accuracyAverage_50_50))
print('RT_accuracyTestList (20% train, 80% test) partition mean: ' + str(RF_accuracyAverage_20_80))
```

RT_accuracyTestList (80% train, 20% test) partition mean: 1.0
RT_accuracyTestList (50% train, 50% test) partition mean: 0.932584269663
RT_accuracyTestList (20% train, 80% test) partition mean: 0.889671361502

In [109]:

```
displayAccuracies('RT', 'Wine Data', RF_accuracyAverage_80_20, RF_accuracyAverage_50_50, RF_accuracyAverage_20_80)
printAccuracies('RT', RF_accuracyTestList_80_20, RF_accuracyTestList_50_50, RF_accuracyTestList_20_80)
```



Accuracy of RT's 3 trials on (80% train, 20% test) partition :[1.0, 1.0, 1.0]

Mean Accuracy of RT on (80% train, 20% test) partition: 1.0

Accuracy of RT's 3 trials on (50% train, 50% test) partition :[0.9550561797752809, 0.8764044943820225, 0.9662921348314607]

Mean Accuracy of RT on (50% train, 50% test) partition: 0.932584269663

Accuracy of RT's 3 trials on (20% train, 80% test) partition:[0.89436619718309862, 0.88732394366197187, 0.88732394366197187]

Mean Accuracy of RT on (20% train, 80% test) partition: 0.889671361502

Results

In [103]:

```
def autolabel(rects):  
    """  
    Attach a text label above each bar displaying its height  
    """  
    for rect in rects:  
        height = rect.get_height()  
        ax.text(rect.get_x() + rect.get_width()/2., 1.05*height,  
                #'%d' % int(height),  
                '%d' % int(height) + '%',  
                ha='center', va='bottom')
```

In [104]:

```
import numpy as np
import matplotlib.pyplot as plt

# data to plot
n_groups = 3
SVM_partitions = (SVM_accuracyAverage_80_20*100, SVM_accuracyAverage_50_50*100,
SVM_accuracyAverage_20_80*100)
DT_partitions = (DT_accuracyAverage_80_20*100, DT_accuracyAverage_50_50*100, DT_
accuracyAverage_20_80*100)
RT_partitions = (RF_accuracyAverage_80_20*100, RF_accuracyAverage_50_50*100, RF_
accuracyAverage_20_80*100)

# create plot
fig, ax = plt.subplots()
index = np.arange(n_groups)
bar_width = .25
opacity = .7

SVM = plt.bar(index, SVM_partitions, bar_width,#align='center',
               alpha=opacity,
               color='r',
               label='SVM')

DT = plt.bar(index + bar_width, DT_partitions, bar_width,#align='center',
               alpha=opacity,
               color='b',
               label='Decision Tree')

RT = plt.bar(index + bar_width + bar_width, RT_partitions, bar_width,#align='cen
ter',
               alpha=opacity,
               color='g',
               label='Random Forest')

plt.xlabel('Partitions')
plt.ylabel('Test Accuracy')
plt.title('Test Accuracies of Classifiers on Wine Data', y=1.15)
plt.xticks(index + bar_width, ('80%train,20%test', '50%train,50%test', '20%train
,80%test'))
#plt.legend()
ax.legend(loc='center left', bbox_to_anchor=(1, 0.5))

autolabel(SVM)
autolabel(DT)
autolabel(RT)

plt.tight_layout()
plt.grid()
plt.show()
```

Test Accuracies of Classifiers on Wine Data

